



Ioana Frîncu

TECHNICAL PRODUCT MANAGER · AI PLATFORMS MLOPS · HUMAN-AI INTERACTION · TECHNICAL AI SAFETY

Brussels, Belgium · Willing to relocate

☎ (+32) 456390234 | ✉ f_ioana@tuta.com | 📱 flifenstein | 🌐 frincui

“Stuff happens. The fun part is figuring out how to fix it ... for the long term.”

Summary

Technical PM with an HCI master's shipping **AI infrastructure products** — from model-serving platforms (KServe, vLLM) at OpenChip to applied research tooling at Microsoft's AI for Good Lab, and now at imec's SiliconAI team building an agentic inference benchmarking system. I'm the team's **“glue worker”**: I turn ambiguous ML problems into shipped products while making chief architects and nonprofit stakeholders actually understand each other. Now looking to transition into **technical AI safety** — returning to evaluation work that started with my MSc thesis on LLM prompt-tuning (cited in Abdulhai et al. 2025 on reducing LLM deception), and building ML projects on the side to sharpen the technical edge. Looking for a seat on a technical AI safety team where I can own **0-to-1 work** alongside people doing the hard research.

Work Experience

R&D Product Manager — AI Team (Prof. Steven Latré)

IMEC

Ghent, Belgium

May 2026 - Ongoing

- Joining imec's SiliconAI applied AI research group as part of the OpenChip AI platform team. Scoping a new benchmarking system for **evaluating LLM behavior under realistic agentic inference workloads**. Bridging production inference infrastructure and applied LLM evaluation research.

AI Platform Product Owner

OPENCHIP

Ghent, Belgium

Nov. 2025 – Apr. 2026 (6 mo)

- Defined product strategy for Booster, a multi-tenant AI inference platform on **KServe and vLLM** targeting external ML engineers deploying LLMs in production
- Operated as central **decision-maker** across leadership, engineering, architecture, and design for a team of 6 engineers — gathered stakeholder input, initiated cross-team discussions, and closed decisions autonomously when escalation wasn't required (e.g., resolving shared vs. unique tenant identity in multi-tenant deploys).
- Shipped a dark-mode **Claude Code prototype** in under two days as a forcing function to resolve beta-scope debates open for weeks; owned the Jira backlog, Agile process, and spokesperson role for the team.

Product Management Intern — AI for Good Lab

MICROSOFT

Remote / Brussels office

Mar. 2025 – Nov. 2025 (9 mo)

- Drove the end-to-end launch of **LINGUA: Expanding Europe's Voices in AI** — a Microsoft European Digital Commitments open call funding open-licensed datasets for underrepresented European languages. Owned program scoping, applicant rubric, eligibility framework, and partner coordination; the program selected **11 projects across 16 languages and 10 countries**. Coordinated stakeholder alignment across **EPFL & ETH Zürich (APERTUS), Mozilla, and the Council of Europe**
- Scoped responsible AI review processes and Azure grant-management workflows; shipped an internal AI storytelling tool that streamlined the lab's research-to-impact reporting across concurrent projects.
- Engaged with Microsoft's **Brussels EU policy office** on regulatory touchpoints for the lab's European initiatives

Young Graduate — Product & Customer Success Lead

TIWPE DIGITAL PUBLISHING

Leuven, Belgium

Oct. 2023 – Mar. 2025 (1 yr 6 mo)

- Delivered a **multi-title feature for Centre France** on a 21-MD budget — coordinated stack migration, UX improvements, release monitoring, and post-release triage for a multi-brand app consolidation.
- Replaced expensive manual QA with **automated E2E testing**; owned platform releases across mobile and web apps and drove compliance-critical technical deadlines ahead of quarter.

Research Intern — NLP / LLM Evaluation

Remote

IRIS.AI

Feb. 2023 – Sep. 2023 (8 mo)

- Ran a human-subject evaluation study on soft vs. hard prompt-tuning techniques for LLM-generated scientific abstracts (MSc thesis, cited in Abdulhai et al. 2025 (Berkeley/Oxford) on reducing LLM deception via multi-turn RL). Designed study protocol, recruited participants, and analyzed qualitative and quantitative measures of output utility across prompting strategies.
- Hands-on work with **LLM fine-tuning, prompt engineering, and human-AI evaluation methodology**. The methodological foundation underpinning my ongoing interest in technical AI safety (evaluations, model behavior, interpretability).

Skills

AI / ML

Model serving, LLM fine-tuning, prompt engineering, RAG, human-AI evaluation; Claude & OpenAI APIs, Hugging Face, Ollama mini-projects (Github)

AI Safety & Responsible AI

LLM evaluation design, BlueDot Technical AI Safety curriculum, AI ethics frameworks (Turing AI Ethics), RAI review processes, EU AI Act literacy.

Product Management

Product strategy, 0-to-1 scoping, roadmap & prioritization, technical specs, competitive teardowns, stakeholder alignment, Jira & Agile ownership.

Engineering & Data

Python, SQL, R, JavaScript & TypeScript, React, TailwindCSS, JupyterLab; Kubernetes and Jenkins (platform familiarity).

UX & Prototyping

Figma, Claude Code, Lovable, InDesign; qualitative and quantitative user-research methods.

Languages

English (C2), Romanian (C2), Dutch (B2), French (A2).

Education

MSc, Human-Computer Interaction & Design — Double Degree

Netherlands & Finland

EIT DIGITAL MASTER SCHOOL

Sep. 2021 – Sep. 2023

- University of Twente (Year 1) + Aalto University (Year 2). Minor in Innovation & Entrepreneurship. Technical coursework: Machine Learning, NLP, Conversational Agents (incl. Advanced Project), Speech Processing, Affective Computing, Product Management, Computer Ethics.
- **Thesis** on LLM prompt-tuning evaluation (Aalto repository; cited in Abdulhai et al. 2025), conducted with Iris.ai.

BSc, Communication Science

Netherlands

UNIVERSITY OF TWENTE

Sep. 2017 – Sep. 2021

- Coursework: digital marketing, UX, persuasive technology, statistics & R, media studies. Board member, Study Association Communiqué (25th board).

Certifications

May. 2026 **BlueDot**, Technical AI Safety

Online

Feb. 2026 **Turing AI Ethics**, Turing College

Online

Dec. 2025 **Microsoft Certified: AI-900**, Microsoft

Online

Apr. 2024 **Product Management Intro**, Nova

Online

Mar. 2022 **Technology Transfer & Innovation Workshop**, European Space Agency

Online

Dec. 2021 **Google IT Support Specialization (scholarship)**, Coursera

Online

Community & Extracurricular

Co-founder

Brussels, Belgium

WOMEN IN AI BELGIUM

Jan. 2026 – Present

- Founding ambassador for the Belgian chapter — organizing community events and knowledge-sharing sessions on applied AI and responsible AI practice.

Board Member — External Affairs

Enschede, Netherlands

STUDY ASSOCIATION COMMUNIQUÉ

Sep. 2019 – Oct. 2020

- Built industry partnerships for Career Fair, Symposium, and company visits; secured €4,000 in sponsorships